

Pushing Four Raspberry Pis to the Memory Wall

Bit-Exact 30B Mixture-of-Experts Inference,
+16% over the Public Record

Daniel Correa Villa

HELLOMATIK.COM

Raspberry Pi Cluster Lab

hellomatik.com/research/hitting-the-memory-wall-30b-moe-raspberry-pi

23 May 2026

Contents

1	Introduction	4
2	Related work	6
3	System design	7
3.1	Hardware	7
3.2	Topology	8
3.3	Software stack	8
4	Methodology	8
4.1	Metrics	8
4.2	Protocol	8
5	Optimisations applied	9
5.1	Framework source patches	9
5.2	Toolchain	9
5.3	Operating-system tuning	9
5.4	Architectural change: dense to MoE	9
6	Results	9
6.1	Final-configuration throughput and stability	9
6.2	Optimisation trajectory	12
6.3	Memory and thermal behaviour	12
7	Stage 9 — TIER 0 kernel network and scheduler tuning	13
7.1	Configuration changes	13
7.2	Measured impact	13
7.3	Two rejected candidates from the same research round	13
7.4	Open work after Stage 9	13
8	Stages 10 and 11 — Bit-exact MoE op-fusion and Rounds 3/4 ceiling investigation	14
8.1	Stage 10 — A2 OP_SILU_MUL fusion (+0.5%)	14
8.2	Stage 11 — Rounds 3 and 4 (Linux papers research and ceiling investigation)	14
8.3	Bottom line after Stage 11	17
9	Stages 12–15 — Hot-kernel inversion, true single-pass fusion, network chunk-size sweep, runtime NIC tuning	17
10	Stage 16 — WFE/SEV inter-step barrier and API robustness hardening	19
10.1	API robustness: uncaught-exception crash on client disconnect	19
10.2	WFE/SEV barrier	19
10.3	Measurement	19
11	Statistical significance of the Stages 8–15 trajectory	20
12	Failed configurations and frameworks	20
13	Discussion	21
13.1	Bottleneck analysis	21
13.2	Mixture-of-Experts as architectural sparse computation	22
13.3	Comparative cost/performance	22
13.4	Limitations	22
13.5	Future work: HALO overlap scheme	22
13.6	Future work: alternative model candidates	23
14	Clean-room decode-rate replication and telemetry interference	23
14.1	Motivation: metric alignment with the public ceiling	23
14.2	Telemetry-bandwidth interference	23
14.3	Apples-to-apples result against the public ceiling	24

14.4 Bit-exact kernel work (T0.1) — partial success, reverted	24
15 Final conclusion	25
15.1 What moved the needle	25
15.2 Operating-system tuning	25
15.3 The memory wall	26
15.4 Future work	26
15.5 Threats to validity	26

Abstract

A 30-billion-parameter Mixture-of-Experts language model (Qwen3-30B-A3B) runs on four Raspberry Pi 5 boards — roughly €500 of CPU-only hardware, with no GPU or NPU — at 15.143 tok/s decode, bit-exact. On *identical* 16 GB silicon this is **+15.2%** over the vanilla `distributed-llama` baseline (13.15 $\rightarrow$ 15.143 tok/s, $n=20$) — a confound-free figure that isolates the source-level and runtime-configuration work reported here. It is also **+16.1%** above the highest publicly documented configuration for this model and hardware class (13.04 tok/s; b4rtaz #255, 8 GB SKU); we show that cross-SKU comparison carries no measurable hardware confound, since vanilla `distributed-llama` on our 16 GB cluster measures 13.15 tok/s — within run-to-run noise of the 8 GB record.

We report the design, optimisation and rigorous benchmarking of a four-node Raspberry Pi 5 (16 GB) cluster for distributed inference of this Mixture-of-Experts (MoE) model (3B active parameters per token, top-8 of 128 experts) over Gigabit Ethernet.

After exhausting framework alternatives infeasible on ARM Linux without GPU/NPU acceleration (EXO requires Apple MLX; `prima.cpp` hangs on ZMQ topology discovery; `llama.cpp` RPC regresses 25 $\times$), we adopt `distributed-llama` v0.16.5 with twelve source-level changes developed in this work — eight framework patches and four bit-exact kernel and op-fusion optimisations — plus persistent runtime kernel tuning. Every optimisation stage is individually validated bit-exact (SHA-256 of the generated token-id sequence against a fixed reference at `seed=42`, `temperature=0`).

The trajectory runs from a 5.70 tok/s Llama-3.1-8B dense baseline, through an 11.40 tok/s Qwen3-MoE baseline, to the 15.143 tok/s decode result above ($n=20$, 95% CI ± 0.097), measured with the exact #255 command (`nTokens=109` on every run); end-to-end sustained serving throughput, prefill included, is 14.449 tok/s.

ARM PMU profiling locates the residual bottleneck in LPDDR4X memory bandwidth (49% backend-stalled cycles, 11.4 GB/s sustained per node of an approximate 17 GB/s ceiling). We catalogue twenty-six unsuccessful configurations, document the constraint-discovery process (strict bit-exact, no kernel rebuild, no model change, no overclock) that narrowed the design space, and find that co-located observability agents (Grafana Alloy, cAdvisor) tax this memory-bound decode by 5.18% while leaving compute-bound prefill unchanged — a clean, same-hardware confirmation that decode is bandwidth-bound.

The cluster reaches the memory wall of the Raspberry Pi 5; the single remaining bit-exact software lever is Async Tensor Parallelism, an estimated 250 LOC integration on the Phase B foundation already shipped in the repository.

Table 1. Results at a glance.

Decode throughput (apples-to-apples vs. the public ceiling)	15.143 tok/s
improvement over b4rtaz #255 (13.04 tok/s, decode)	+16.1%
improvement over vanilla on identical 16 GB silicon (confound-free)	+15.2%
Sustained serving throughput (prefill included)	14.449 tok/s
Time-to-first-token (TTFT)	557 ms
Per-node DRAM bandwidth (sustained / vendor ceiling)	11.4 / 17 GB/s
Hardware / cost	4 $\times$ Pi 5 16GB / $\sim$ €500
Telemetry-off gain on decode (paired A/B)	+5.18%
Constraints honoured	bit-exact, no overclock, no model change

1 Introduction

Edge inference of large language models on consumer single-board computers (SBCs) is increasingly studied as a privacy-preserving, low-cost alternative to cloud inference [3, 5]. The Raspberry Pi 5 is the most widely deployed ARM-based SBC with enough RAM (16 GB) to host 8B-parameter quantised models, and its Ethernet I/O allows small clusters to be formed.

The absence of a usable GPU — the VideoCore VII lacks general compute shaders — restricts inference to the CPU, where execution is memory-bandwidth-bound.

Investigation timeline at a glance.

The investigation comprised three intertwined tracks:

- (i) **Architectural exploration** (framework, model, quantisation): converged on `distributed-llama` + Qwen3-30B-A3B Q40, establishing the operational baseline of 11.40 tok/s.

- (ii) **Production-branch source work** (Stages 1–8, 7.18 → 13.720 tok/s): driven mainly by the Stage 4 switch from Llama 3.1 8B to Qwen3-MoE (+59%), plus the F0/F1/F2 **distributed-llama** patches.
- (iii) **Bit-exact refinement** (Stages 9–15, 13.720 → 14.449 tok/s sustained; 15.143 tok/s decode): documented in detail here.

Five hard constraints, discovered progressively and adopted as design rules, frame the entire work:

- *Bit-exact output* — SHA-256 of the first 100 generated token-ids matches a fixed pre-stage reference (`seed=42`, `temperature=0`).
- *No kernel rebuild* — stock Pi OS kernel 6.12.75.
- *No model change* — Qwen3-30B-A3B Q40 is fixed.
- *No CPU overclock* — silicon stays at the nominal 2.4 GHz.
- *No quality reduction* — top- k unchanged at 8, no Q3 down-quantisation, no expert pruning.

These constraints rule out most speedup techniques in the recent literature [11, 15], but they make the surviving optimisations production-deployable with no risk of quality regression.

Contributions of this report.

1. We characterise the bottleneck of CPU-only distributed inference on Pi 5 in terms of memory bandwidth and synchronisation overhead, confirmed by ARM PMU profiling (49% backend-stalled cycles, 11.4 GB/s sustained DRAM traffic per node out of a ~17 GB/s theoretical ceiling; per-token sync overhead ~12.3 ms = 17.7% of token time).
2. We develop and apply **twelve** source-level changes to the **distributed-llama** framework (eight framework patches plus four bit-exact kernel and op-fusion optimisations) that fix critical bugs (uint8 overflow, JSON crash, header truncation), enable optimisation flags previously inaccessible, introduce a bit-exact `OP_SILU_MUL` kernel fusion (Stage 10) and a true single-pass version of the same kernel (Stage 13), invert a long-standing software-prefetch micro-optimisation in the matmul inner loop (Stage 12, the HW prefetcher of the Cortex-A76 turns out to be more efficient than the manual hints it was given), tune `MAX_CHUNK_SIZE` for the `writeMany/readMany` all-reduce loop (Stage 14, 4-fold reduction in `send()/recv()` syscall count), and ship a foundation for tile-level overlap of computation and communication for a future sprint.
3. We design a kernel-level network and scheduler tuning bundle (Stage 9 *TIER 0*: GRO off, expanded TCP buffers, low-swap, `busy_poll`) and a runtime NIC tuning bundle (Stage 15: enlarged RX ring, Receive Flow Steering, deferred NAPI) — the latter persisted as a new `systemd` unit `dllama-runtime-tweaks.service` — and quantify the throughput contribution of each.
4. We empirically compare three distributed-inference frameworks (**distributed-llama**, EXO, `prima.cpp`) and document failure modes for **twenty-six** model, framework and optimisation configurations, including the hypotheses rejected post hoc in the most recent investigation round (e.g. ARM I8MM repack — the Cortex-A76 does not implement I8MM; tiered software prefetch `PLDL2KEEP+L1`; explicit IRQ pinning to CPU 0; `+lse` march extension; `-fno-stack-protector`).
5. We report rigorous benchmarks ($n=20$ paired runs after two discarded warm-up runs, fixed prompt, `temperature=0`, 95% confidence intervals, p50/p90/p99 percentiles) for every stage of the trajectory, with **bit-exact output verification** at `--seed 42` via SHA-256 hashing of the generated token-id sequence against a fixed reference. Every improvement claim in this paper is validated bit-exact unless explicitly marked otherwise.

2 Related work




Figure 1. Same model and hardware class (Qwen3-30B-A3B Q40 on 4×Raspberry Pi 5), *decode* metric: this work versus the highest publicly documented result (b4rtaz #255), a bit-exact +16.1% improvement (15.143 vs. 13.04 tok/s, $n=20$, telemetry off).

The state of the art for distributed CPU LLM inference falls into three approaches:

- **Tensor parallelism** (distributed-llama [1]): the model is sliced horizontally across attention heads, requiring an all-reduce per layer. Used in this work.
- **Pipeline parallelism** (llama.cpp RPC, prima.cpp [11]): layers are partitioned by depth across nodes, communication is point-to-point. Geerling [3] reports a 25× regression vs. single-node for llama.cpp RPC, confirming the network-overhead-dominated regime.
- **Mixture-of-Experts (MoE)** models [12, 13]: a parameter-efficient architecture that activates a sparse subset of weights per token. Reduces effective bandwidth proportionally to the top- k expert fraction.

Several comparable Pi 5 clusters appear in the literature. Tadych [1, 2] reports 13.04 tok/s on 4×Pi 5 8GB with Qwen3-30B-A3B; the Seed Studio documentation [4] reports 6.06 tok/s with DeepSeek-R1-Distill-Llama-8B on the same hardware; and the Stratosphere Laboratory [5] provides single-node Pi 5 baselines. Table 2 positions this work against the most directly comparable published systems.

Table 2. Positioning of this work against the most directly comparable published distributed-inference systems for the same model class (Qwen3-30B-A3B Q40 or equivalent MoE) on Pi 5-class hardware. Throughput figures are as reported by each source on its own hardware; the right column gives the closest comparison that can be made against *our* hardware (4×Pi 5 16 GB cluster) when the published source benchmarked on different hardware.

System	Framework / approach	ap- tok/s	Quality	Notes
This work (Stage 17)	distributed-llama + 12 patches + kernel/NIC tuning	15.143	bit-exact (SHA-256)	4×Pi 5 16 GB, this paper; decode metric, telemetry off; +16.1% vs Tadych ceiling (sustained 14.449 tok/s)
Tadych #255 [2]	distributed-llama v0.16, default	13.04	bit-exact	4×Pi 5 8 GB, vanilla upstream (decode)
Sseed Studio [4]	distributed-llama	6.06	bit-exact	4×Pi 5 8 GB, DeepSeek-R1-Distill-Llama-8B
ByteShape [20]	llama.cpp , single node	8.03	bit-exact	1×Pi 5 16 GB, Q3_K_S 2.70 bpw
Geerling [21]	llama.cpp , single node	7.99	bit-exact	1×Pi 5 16 GB, comparable quant
Geerling RPC [3]	llama.cpp RPC distributed	—	bit-exact	Reported 25× regression vs. single-node on Pi clusters
prima.cpp [11]	pipelined ring + mmap	n/a	bit-exact	ZMQ topology fails to bootstrap on Pi 5 in our testing
HALO [15]	distributed-llama + 3,500 LOC overlap	—	bit-exact	1.30× on reliable LAN (their hardware); source not yet public, speedup not independently verified here
EXO [14]	MLX backend	n/a	—	Requires Apple MLX; not buildable on ARM Linux

For methodology we follow the MLPerf Inference v5.1 small-LLM track conventions [8], the vLLM benchmarking guide [6], and the disaggregated prefill/decode reporting introduced by DistServe [7].

3 System design

3.1 Hardware

Table 3. Cluster hardware inventory.

Node	Role	LAN IP	Storage
rpi-1005	root + serving	192.168.1.74	NVMe 457 GB
rpi-1006	worker	192.168.1.77	NVMe 457 GB
rpi-1007	worker	192.168.1.75	NVMe 457 GB
rpi-1008	worker	192.168.1.76	NVMe 457 GB

Per-node specifications.

- SoC: Broadcom BCM2712 (4× Arm Cortex-A76 @ 2.4 GHz, ARMv8.2-A)
- Memory: 16 GB LPDDR4X @ 4267 MT/s, theoretical bandwidth $\approx$ 17 GB/s
- Storage: NVMe SSD via PCIe Gen 2 ($\approx$ 700 MB/s read)
- Network: integrated Gigabit Ethernet; measured intra-cluster latency 0.226 ms
- Operating system: Debian 13 *trixie*, kernel 6.12.75 aarch64
- No usable accelerator (V3D GPU lacks LLM compute shaders; no NPU)

3.2 Topology




Figure 2. Cluster topology. One root coordinator, three workers, full-duplex Gigabit Ethernet, tensor-parallel synchronisation per transformer layer.

3.3 Software stack

- **Inference engine:** distributed-llama v0.16.5 with eight framework patches (§5); four further bit-exact kernel and op-fusion optimisations follow in Stages 10–14.
- **Model:** Qwen3-30B-A3B Q40 (30B total parameters, 128 experts, top- $k=8$, ~ 3 B active per token).
- **HTTP front-end:** custom Python proxy (Flask + waitress) on 127.0.0.1:8000, sanitising OpenAI-strict client requests.
- **Allocator:** jemalloc 2 preloaded via LD_PRELOAD.
- **Client:** Hermes Agent v0.13.0 for autonomous-agent integration testing.

4 Methodology

We follow the MLPerf Inference v5.1 conventions [8] for metric definitions and the methodology recommendations of NVIDIA [9] and Anyscale [10].

4.1 Metrics

- **Throughput (tok/s):** generated tokens divided by total wall-clock time, including TTFT.
- **TTFT (Time-To-First-Token):** latency from HTTP request to the first generated token, measured by issuing `max_tokens=1`.
- **Prefill rate:** prompt tokens processed per second. Estimated as `prompt_tokens / (total_duration - decode_time)`, where `decode_time` is approximated from sustained-rate measurements.
- **Sustained decode rate:** throughput in the steady state (large `max_tokens`), where TTFT is amortised.
- **Percentile latencies (p50, p90, p99):** ordered statistics over n runs.

4.2 Protocol

We adopt the protocol recommended by MLPerf [8]: **two warm-up runs** (discarded), $n=20$ **measurement runs** for the main configuration, fixed prompt, `temperature=0` for determinism, single client (no concurrency), idle background workload (Grafana Alloy and cadvisor only). Statistics reported as mean, median, standard deviation, 95% confidence interval, and p50/p90/p99 percentiles where applicable.

5 Optimisations applied

5.1 Framework source patches

Table 4. Source-level patches applied to distributed-llama v0.16.5.

#	Patch	Location	Rationale
1	NnByte→NnUInt for nBatches	nn-cpu-ops.hpp:14	uint8_t overflow: setting nbatches=256 silently became 0 (256 mod 256), tripping an embedding-layer assertion at runtime.
2	Force finish_reason = "stop"/"length"	dllama-api.cpp:547	Empty finish_reason caused strict OpenAI clients (Hermes) to enter infinite retry loops.
3	try/catch around json::parse	dllama-api.cpp:83	nlohmann::json threw uncaught exceptions on malformed bodies, abort-trapping the daemon (SIGABRT).
4	headerData.append(buffer, bytesRead)	dllama-api.cpp:123	Default append(const char*) truncated at the first \0 byte, corrupting bodies that contained UTF-8 multi-byte sequences with NUL.
5	New CLI flag --nbatches	app.cpp:130	nBatches was hard-coded to 32; the flag now permits configuration once patch #1 is applied.
6	posix_memalign(64, n) for pipes	nn-executor.cpp:18	Default new[] alignment is 16 B on ARM64; cache-line (64 B) alignment is required for vectorised NEON access.
7	TCP SO_RCVBUF/SO_SNDBUF=8 MB	nn-network.cpp:60	Default 208 KiB buffers caused write-blocking under burst sync at end-of-layer.
8	buffer.reserve(64 KB)	dllama-api.cpp:475	The streaming response buffer doubled repeatedly during long outputs, causing reallocations on the hot path.

5.2 Toolchain

Recompiled with CXXFLAGS="-O3 -flto -ffast-math -funroll-loops" on each node. LD_PRELOAD=libjemalloc.so.2 added to all systemd units with MALLOC_CONF="narenas:4,tcache:true,dirty_decay_ms:30000".

5.3 Operating-system tuning

- CPU governor performance on all 4 nodes (systemd unit at boot).
- TCP BBR congestion control; net.ipv4.tcp_low_latency=1; net.core.netdev_budget=600.
- vm.overcommit_memory=1, vm.swappiness=1 (lowered from the Pi OS default of 60 in Stage 9 to keep mlock-ed weights resident; verified live at 1).
- NVMe scheduler none, read_ahead_kb=2048.
- Swap on NVMe: 16 GB on root, 4 GB on workers.
- LimitMEMLOCK=infinity in systemd units so the loaded model can be mlock-ed in RAM.
- Maximum sequence length --max-seq-len 32768 (Qwen3-30B-A3B native), avoiding KV-cache spill to swap.
- Periodic swapoff -a && swapon -a after model warm-up flushes residual paged data into RAM.

5.4 Architectural change: dense to MoE

The most impactful optimisation was the migration from **Llama 3.1 8B** (dense, ~5 GB Q40 weights, all parameters active per token) to **Qwen3-30B-A3B** (MoE, 128 experts, top- $k=8$, ~3 GB active weights per token). Bandwidth-bound throughput improved by 59% (7.18 → 11.4 tok/s) with no other change.

6 Results

6.1 Final-configuration throughput and stability

The final result, on the *decode* metric that the public ceiling reports, is **15.143 tok/s** (Stage 17 clean-room replication of b4rtaz #255, $n=20$, 95% CI ± 0.097 , telemetry off; Section 14), +**16.1%** above the 13.04 tok/s ceiling. The end-to-end *sustained* serving throughput (250-token responses, prefill included) is **14.449 tok/s** (cold-cluster, $n=20$ paired runs, 95% CI ± 0.038 , best individual run 14.557 tok/s), reached at Stage 15 through the optimisation trajectory of Section 6.2 (Table 5); the two are consistent, as 15.143 tok/s of decode amortises to ≈ 14.5 tok/s once the ~ 557 ms prefill is included.

The detailed run-to-run distribution, TTFT and response-length characterisation that follow were captured at the *Stage 6 intermediate snapshot* (12.708 tok/s sustained); they document the cluster’s run-to-run

stability and latency profile, which the later Stage 9–15 optimisations raised in throughput without changing the qualitative latency behaviour. Each distribution is reported at the snapshot where it was measured rather than restated at the final configuration.

Table 5. Statistical summary of sustained throughput at the final configuration (Stage 15), Qwen3-30B-A3B on 4×Pi 5.

Statistic	Value
Mean (with runtime tweaks)	14.449 tok/s
95% CI ($\pm$)	0.038 tok/s
Best individual run	14.557 tok/s
Mean (without runtime tweaks)	14.167 tok/s
95% CI ($\pm$)	0.043 tok/s
n (paired runs)	20

Per-run throughput distribution (Stage 6 snapshot).

Table 6. Statistical summary of throughput at the Stage 6 intermediate snapshot, Qwen3-30B-A3B on 4×Pi 5.

Statistic	Value
Mean	12.708 tok/s
Median (p50)	12.707 tok/s
Standard deviation	0.066 tok/s
Min	12.585 tok/s
Max (p100)	12.852 tok/s
p90	12.812 tok/s
p99	12.852 tok/s
95% CI ($\pm$)	0.029 tok/s
Coefficient of variation	0.52%




Figure 3. Per-run throughput across 20 measurement runs at the Stage 6 intermediate snapshot (warm-up runs excluded). Coefficient of variation 0.52%.

Time-To-First-Token (Stage 6 snapshot).

Table 7. TTFT statistics over $n=10$ runs with `max_tokens=1`.

Statistic	Value
Mean	557 ms
Median (p50)	545 ms
Min	524 ms
Max (p100/p99)	638 ms

Throughput scaling with response length (Stage 6 snapshot).




Figure 4. Sustained throughput as a function of response length at the Stage 6 intermediate snapshot. Asymptotic throughput converges near 12.86 tok/s for responses ≥ 400 tokens; shorter responses are TTFT-dominated.

Prefill rate vs. prompt size (Stage 6 snapshot).

Table 8. Prefill rate (estimated) for increasing prompt lengths.

Prompt tokens	Total time (s)	Est. prefill time (s)	Prefill rate (tok/s)
57	3.66	3.02	18.9
417	23.97	23.33	17.9
2017	129.07	128.43	15.7

Prefill rate degrades slightly with prompt size (KV-cache growth and the attention $O(N^2)$ cost in the longer-prompt regime) but remains >15 tok/s up to 2K context. For 20K-token prompts (e.g. Hermes Agent system prompts), the inferred prefill time would exceed 20 minutes — the practical upper limit of usability for this configuration on full agent workloads.

6.2 Optimisation trajectory




Figure 5. Optimisation trajectory and final clean-room result (Qwen3-30B-A3B Q40, 4×Raspberry Pi 5). Bars 1–8 report sustained throughput; the dashed line marks the public ceiling (13.04 tok/s, b4rtaz #255). The Stage 14–15 configuration reaches 14.449 tok/s sustained ($n=20$, 95% CI ± 0.038); the final green bar is the Stage 17 clean-room #255 replication on the *decode* metric — the quantity directly comparable to the ceiling — at **15.143 tok/s** ($n=20$, 95% CI ± 0.097 , telemetry off), +16.1% above the ceiling. The Llama 3.1 8B dense bar is the abandoned starting point, shown for context.

6.3 Memory and thermal behaviour




Figure 6. Memory utilisation per node during sustained inference. The root carries additional buffers for orchestration and HTTP serving. Workers retain >9 GB headroom for further KV-cache growth.

Table 9. Sustained-load thermal measurements. Throttle threshold is 85°C.

Node	Temperature	Throttling flags
rpi-1005 (root)	54.3 °C	0x0
rpi-1006 (worker)	54.3 °C	0x0
rpi-1007 (worker)	55.4 °C	0x0
rpi-1008 (worker)	56.0 °C	0x0

7 Stage 9 — TIER 0 kernel network and scheduler tuning

After the initial publication of this report at 12.71 tok/s and the subsequent Stages 6–8 documented in the project README (which raised sustained throughput to 13.720 tok/s), a structured research round identified a further set of OS-level network and memory tunings that yielded a measurable improvement with no source-code changes.

7.1 Configuration changes

Applied to all four Pi 5 nodes, persisted via `/etc/sysctl.d/99-llama.conf` and an `eth0-tuning.service` systemd unit:

Setting	Before	After
<code>ethtool -K eth0 gro</code>	on	off
<code>net.core.rmem_max</code>	212 992	8 388 608
<code>net.core.wmem_max</code>	212 992	8 388 608
<code>net.ipv4.tcp_rmem (mid/max)</code>	131 072 / 6 291 456	87 380 / 8 388 608
<code>net.ipv4.tcp_wmem (mid/max)</code>	16 384 / 4 194 304	65 536 / 8 388 608
<code>vm.swappiness</code>	60	1

Table 10. Stage 9 sysctl and `ethtool` changes.

Rationale. GRO (Generic Receive Offload) coalesces incoming packets, adding 50–200 us of intentional latency; for the 510 kB sync bursts of the all-reduce step on a 1 GbE LAN this is pure overhead. Raising `rmem_max/wmem_max` from the 256 kB default lets the TCP receive/send windows grow past the burst size. The 16 kB `tcp_wmem` middle value was a hard bottleneck for the worker→root activation send.

7.2 Measured impact

We re-baselined (after the Phase B foundation work) at **13.720 ± 0.050 tok/s** ($n=20$). The post-Stage 9 measurement ($n=20$, paired, 95% CI) is:

Metric	Pre-Stage 9	Post-Stage 9
Mean	13.720 tok/s	14.011 tok/s
Stdev	0.050	0.116
95% CI	±0.022	±0.051
Δ vs. Stage 8 baseline	—	+2.12%
Δ vs. public ceiling (13.04)	—	+7.45%

7.3 Two rejected candidates from the same research round

The same research round surfaced two further claims, both investigated and rejected:

1. *llamafile_sgemm guard fix* (upstream issue #284, claimed +5–40%): this gain had *already* been captured in our base commit 92a20e2 ‘‘perf: enable `llamafile_sgemm` for single-token decode’’ (2026-05-16), two days before this research round. Cannot be captured twice.
2. *ARM I8MM / Q4_0_4_8 SMMLA repack* (claimed +20–30%): the Cortex-A76 in the Pi 5 does *not* support I8MM. Verified directly: `grep -c i8mm /proc/cpuinfo` returns 0 on all four nodes. Confirmed via llama.cpp issue #10662. I8MM is an A78+ feature; the original claim was a factual error in the surveyed material and is documented here as a cautionary tale.

7.4 Open work after Stage 9

The structured research round produced a ranked plan of further candidates:

- **Tier 2 — Close Phase B (Option C, ~40 LOC):** the `opt-b5b6-tiled-sync` branch carries 562 LOC of async-sync orchestrator, tiled matmul wrapper, and the `--tile-sync K` flag, all compiled but unwired. Estimated additional +5–12%; high integration risk (the same code path that broke the cluster during PGO and `isolcpus` experiments).

- **Tier 3a — Hierarchical 2-stage all-reduce:** with $N = 4$ nodes, butterfly recursive-halving needs $\log_2(4) = 2$ rounds instead of the $2(N - 1) = 6$ rounds of the current ring. Estimated +10–13%, ~80–120 LOC.
- **Tier 3b — Pre-attention expert prediction** (arXiv:2511.10676): specific to Qwen3-30B-A3B, reported 94.69% top- k accuracy. Estimated +8–15%, ~250 LOC, low risk (fallback to on-demand load).
- **Tier 4 — Fused `ffn_up` + `ffn_gate` + `silu` MoE op** from `ik_llama.cpp` PR #229; estimated +8–12%, ~400 LOC.
- **Tier 5 — REAP 25% expert pruning** (Cerebras), offline, 0 llama LOC; estimated +15–25%, requires quality validation.

Bottom line at Stage 9 (mid-investigation snapshot). At this point sustained throughput was **14.011 tok/s** (± 0.051 at 95% CI, $n=20$), +7.45% above the 13.04 ceiling. This was the value at Stage 9; Stages 10–15 later raised sustained throughput to 14.449 tok/s, and the Stage 17 clean-room replication establishes the final decode result of **15.143 tok/s** (+16.1% vs ceiling, Section 14).

8 Stages 10 and 11 — Bit-exact MoE op-fusion and Rounds 3/4 ceiling investigation

After the Stage 9 result (14.011 tok/s), we carried out further bit-exact optimisation work over an extended session. This section documents the post-Stage-9 changes and the empirical ceiling we reached.

8.1 Stage 10 — A2 OP_SILU_MUL fusion (+0.5%)

The Qwen3-MoE per-expert FFN ops are w_1 -matmul $\rightarrow w_3$ -matmul $\rightarrow \text{silu}(w_1) \rightarrow w_1 * w_3$. We fused the last two element-wise ops into a single new `OP_SILU_MUL` with one barrier instead of two, eliminating 28 layers $\times$ 8 experts = 224 barriers per token. The kernel chains the existing `silu_F32` and `mul_F32` kernels verbatim, so the output is byte-identical to the previous two-op sequence (validated with a deterministic prompt at `temperature=0`).

Measurement ($n=20$): 14.081 ± 0.055 tok/s vs 14.011 ± 0.051 baseline. The real gain is +0.50% — modest, but bit-exact and persistent. Committed at SHA 106eab3.

8.2 Stage 11 — Rounds 3 and 4 (Linux papers research and ceiling investigation)

We carried out four parallel research threads:

1. Linux Plumbers / SOSP / OSDI / EuroSys / USENIX ATC papers 2024–2026.
2. eBPF / XDP / AF_XDP / io_uring / kernel-bypass.
3. Linux scheduler (EEVDF in 6.6+) tunings.
4. ARM-specific kernel cache / prefetcher / NUMA / memory bandwidth.

We applied and measured ~15 distinct configurations. **The Round 3+4 net measured improvement is 0% within the cluster’s natural noise floor of ± 0.15 tok/s.**

8.2.1 Literature survey: techniques evaluated and rejected

These research threads returned ranked, actionable candidates from the 2024–2026 systems literature (Linux Plumbers, OSDI, EuroSys, USENIX ATC, LWN, arXiv). Table 11 records each technique, its source, whether it applied to a stock Pi OS 6.12 / Cortex-A76 / 1 GbE cluster under our constraints, and the empirical outcome.

The recurring reason for inapplicability is structural: the highest-leverage modern techniques require either a newer kernel, a hardware feature the A76 silicon lacks, or a reboot or kernel rebuild we had ruled out — which is precisely what makes the surviving bit-exact wins deployable on stock hardware.

Table 11. Techniques surveyed during the research round, with applicability to a stock Pi 5 / kernel 6.12 / 1 GbE cluster under the bit-exact, no-reboot, no-quality-loss constraints.

Technique	Source	Applicable here?	Outcome
IRQ suspension (<code>irq_suspend_timeout</code>) MADV_HUGEPAGE	LWN 1008399 Linux mm docs	No — needs kernel 6.13+ No — THP not compiled in Pi OS	Pi OS on 6.12 Tested, no-op
AF_XDP zero-copy	kernel networking docs	No — <code>macb</code> lacks native XDP	RFC only
<code>io_uring</code> SQPOLL DEFER_TASKRUN SO_BUSY_POLL + RPS/RFS Threaded NAPI NUMA <code>fake=4</code> + local alloc MPAM cache partitioning	+ liburing LWN 540281 kernel NAPI docs Linux NUMA docs docs.kernel.org	Net-negative on 4 cores Already in TIER 0 Tested −3.9% Requires reboot Kernel 6.19+ and A76 lacks HW	Skipped Applied Reverted Excluded Future kernel
Per-task EEVDF <code>sched_runtime</code> <code>ik_llama</code> pair-row matmul ILP Jumbo frames MTU 9000	LWN 969062 ik_llama PR <code>macb</code> driver	Applied (4 ms slice) OoO core already optimal <code>bcmgenet</code> EBUSY on live link	No measurable gain Tested, reverted Documented dead-end
Phase B <code>--tile-sync 4</code> (no Option C) Tile-level compute/comms overlap	this work FlashOverlap, arXiv:2504.19519	Tested −8% (unwired) Yes — future sprint	Reverted; foundation in repo Phase B Option C

8.2.2 Phase B foundation cherry-picked

We cherry-picked four commits from the archived `opt-b5b6-tiled-sync` branch onto `pi5-cluster:ce7b96e` (the `--tile-sync K` flag), `6ebd00f` (the `writeMany+readMany` fix), `469101e` (the `async sync orchestrator`), and `2aba914` (the `matmul_Q80_Q40_F32_tiled` wrapper).

With `--tile-sync 0` (the default) the cluster behaves identically (measured 14.092 ± 0.044 tok/s). With `--tile-sync 4` we measured a −8% regression: the wrapper subdivides the syncs but does not drive the matmul to emit per-tile signals (the Option C wiring, ~40 LOC, remains outstanding). The foundation is now in the repo for a future sprint.

8.2.3 EEVDF per-task slice

We called `sched_setattr()` at the start of `executorWorkerLoop` to request a 4 ms slice (the default is ~700 us under EEVDF in kernel 6.12), and confirmed via `/proc/(pid)/sched` that workers report `se.slice = 4000000`. The result, 14.061 ± 0.048 tok/s, shows no measurable gain. The likely reason: with 4 pinned workers, the performance governor, and minimal system load, the default scheduling cadence was already near-optimal.

8.2.4 Sysctl bundles ($R3 + R4$)

Table 12. Round 3 and Round 4 sysctls (persisted, defensive).

File	Key	Value
99-dllama-r3.conf	<code>kernel.sched_autogroup_enabled</code>	0
	<code>net.ipv4.tcp_autocorking</code>	0
	<code>net.ipv4.tcp_slow_start_after_idle</code>	0
	<code>net.ipv4.tcp_no_metrics_save</code>	1
99-dllama-r4.conf	<code>net.ipv4.tcp_thin_linear_timeouts</code>	1
	<code>net.core.netdev_max_backlog</code>	8192
	<code>net.core.rmem_max</code>	33 554 432
	<code>net.core.wmem_max</code>	33 554 432
	<code>net.ipv4.tcp_rmem</code>	4096 87380 33 554 432
	<code>net.ipv4.tcp_wmem</code>	4096 65536 33 554 432
	<code>vm.dirty_bytes</code>	16 777 216
	<code>vm.dirty_background_bytes</code>	4 194 304

Where the time goes: DRAM-bandwidth-bound (49% backend-stalled cycles)




Figure 7. Per-token wall-clock breakdown from ARM PMU profiling (`perf stat`) and `strace`. The workload is DRAM-bandwidth-bound (49% backend-stalled cycles, 11.4 GB/s of ~ 17 GB/s per node): the Q40 matmul dominates and the synchronisation barrier is the only software-addressable slack — the target of the Stage 16 WFE/SEV change.

ARM PMU profiling (`perf stat -e l2d_cache_refill,l2d_cache_wb,bus_access,stalled-cycles-backend,cpu-cycles,instructions` for 60 s during sustained inference):

Table 13. PMU-measured stalls and DRAM traffic during inference (n=4 nodes).

Metric	Root (rpi-1005)	Worker (rpi-1006)
stalled-cycles-backend (% cycles)	49.25%	47.69%
stalled-cycles-frontend (% cycles)	3.83%	—
DRAM read bandwidth	2.74 GB/s	2.61 GB/s
DRAM write bandwidth	8.66 GB/s	8.47 GB/s
Per-node DRAM total	11.40 GB/s	11.08 GB/s
IPC	1.74	—
dTLB miss rate	0.08%	—
iTLB miss rate	0.01%	—
branch-misses	0.07%	—

The 49% backend-stall rate confirms the cluster is **DRAM-bandwidth-bound**: the CPUs wait on memory roughly half the time. The $3.16\times$ write-to-read ratio reflects the multi-pass intermediate-buffer pattern in the MoE FFN, of which the A2 op-fusion (Stage 10) already cuts 224 barriers/token. The TLB miss rates ($<0.1\%$) and branch-mispredict rates ($<0.1\%$) confirm that the existing 16 KB-page configuration is already optimal — transparent hugepages would not help (and are not compiled into Pi OS 6.12 in any case).

On the compute side, `objdump` disassembly of the hot matmul shows 322 `udot/sdot` NEON dot-product instructions in the inner loop at an IPC of 1.74: the kernel is already vectorised to the ARMv8.2-A dot-product ISA limit. This is why every compute-side experiment in Table 11 returns zero gain — the cores sit idle waiting on DRAM rather than starved of issue slots.

8.2.6 Cluster state and benchmark after Stage 11

Table 14. Stage 11 benchmark (mid-investigation snapshot), cold-cluster, $n=20$ paired runs. Superseded by Stages 12–15 (14.449 tok/s sustained) and the Stage 17 clean-room decode result (15.143 tok/s).

Metric	Value
Mean throughput	14.046 tok/s
Median	14.024 tok/s
Standard deviation	0.111 tok/s (CV 0.79%)
95% CI	± 0.049
Δ vs. Stage 1–8 baseline (13.720)	+0.326 (+2.38%)
Δ vs. public ceiling 13.04 (b4rtaz #255)	+1.006 (+7.72%)

8.3 Bottom line after Stage 11

After Stage 11 the cluster sustained **14.046 tok/s** (± 0.049 at 95% CI, $n=20$, cold-cluster), **+7.72% above the publicly documented ceiling** of 13.04 tok/s. The cumulative defensive state (persisted sysctls, EEVDF tuning, the Phase B foundation in the repo, and +0.5% from the Stage 10 fusion) represented a maintainable production baseline at this manuscript’s mid-investigation snapshot.

9 Stages 12–15 — Hot-kernel inversion, true single-pass fusion, network chunk-size sweep, runtime NIC tuning

A subsequent rigorous-bench session in May 2026 added four cumulative bit-exact commits to the production branch, lifting sustained throughput to **14.449 tok/s** (best individual run 14.557, $n=20$, 95% CI ± 0.038) — a further **+2.87% over Stage 11** and **+10.81% over the public ceiling**. Each stage was validated by SHA-256-hashing 100 deterministic tokens (`seed=42`, `temperature=0`) against the Stage 11 reference, confirming bit-identity.

Stage 12 — Remove software prefetch in `matmul_Q80_Q40_F32` (commit 64ef787).

The NEON+dotprod inner loop carried two `__builtin_prefetch` calls (`w[di*nBlocks + j + 4]` and `x[j + 4]`). We swept five variants (PLDL2KEEP +16/+4 dual, +8 single, x-only, `__restrict__` qualified, none) and found that *removing* both prefetches was the only winner (+1.05% over baseline).

The Cortex-A76’s hardware prefetcher (stride detection on sequential `w` access) is more efficient than the manual hints, which compete for issue slots in the 4-wide decoder. Stage 12 therefore reverses a long-standing micro-optimisation that proves counter-productive on this microarchitecture.

Stage 13 — True single-pass `silu_mul_F32` fusion (commit 1af175c).

Stage 10’s `OP_SILU_MUL` eliminated the inter-op barrier, but the kernel was still a thin two-call wrapper (`silu_F32` then `mul_F32`) that kept the intermediate result resident in memory between passes: 3 loads + 2 stores per element.

We rewrote it as a single NEON loop that holds the values in registers throughout: 2 loads + 1 store per element (Algorithm 1), a $\sim 33\%$ reduction in memory ops for this kernel. The measured gain is +1.5% on cluster throughput, bit-exactness preserved (the arithmetic sequence is identical; only the intermediate writeback is removed).

Algorithm 1. Stage 13 single-pass `silu_mul_F32`: bit-exact NEON fusion.

Input: output buffer $y[0..n-1]$ (in-place), multiplier buffer $m[0..n-1]$, thread index t , thread count T
Output: $y[i] \leftarrow \text{silu}(y[i]) \cdot m[i]$ for all i in this thread's split
Invariant: arithmetic sequence (`vrecpeq_f32` + one Newton–Raphson + multiply) is byte-identical to the prior two-pass implementation; only the intermediate writeback to $y[i]$ between the two passes is removed.

```
1:  $(start, end) \leftarrow \text{SPLIT\_THREADS}(n, T, t)$ 
2: for  $i \leftarrow start$  to  $end - 4$  step 4 do           // NEON inner loop, 4 lanes
3:    $x \leftarrow \text{vld1q\_f32}(y + i)$                      // 1 vector load
4:    $mv \leftarrow \text{vld1q\_f32}(m + i)$                    // 1 vector load
5:    $e \leftarrow \text{expf\_neon}(-x)$ 
6:    $d \leftarrow 1 + e$ 
7:    $r \leftarrow \text{vrecpeq\_f32}(d)$                        // reciprocal estimate
8:    $r \leftarrow r \cdot (2 - d \cdot r)$                  // 1 Newton–Raphson iteration
9:    $silu \leftarrow x \cdot r$ 
10:   $out \leftarrow silu \cdot mv$ 
11:   $\text{vst1q\_f32}(y + i, out)$                            // 1 vector store; no intermediate write
12: end for
13: for remaining  $i < end$  do                           // scalar tail
14:    $y[i] \leftarrow (y[i] / (1 + e^{-y[i]})) \cdot m[i]$ 
15: end for
```

Stage 14 — MAX_CHUNK_SIZE sweep, 4 KB $\rightarrow$ 16 KB (commit f8ed9d2).

The `writeMany/readMany` loop in `nn-network.cpp` capped each `send()/recv()` syscall at 4 KB. With the 8–32 MB TCP buffers tuned in Stage 9, this generates $4\times$ more syscalls per slice than necessary. We swept four candidate sizes ($n=20$ each):

Table 15. MAX_CHUNK_SIZE sweep (Stage 14), $n=20$ each.

Chunk size	Mean tok/s	Δ vs 4 KB baseline
4 KB (baseline)	~ 14.27	—
8 KB	14.449	+1.25%
16 KB (selected)	14.449	+1.25%
32 KB	14.346	+0.53%
64 KB	14.40	+0.91%

8 KB and 16 KB are statistically indistinguishable (both 14.449 tok/s, within the ± 0.038 CI) and both sit at the sweet spot — aligned with typical per-worker Pi 5 L1d working sets and below the kernel’s TCP segment-coalescing thresholds. We selected 16 KB for headroom; beyond it (32–64 KB) the larger reads begin to fragment across NIC interrupts and give back part of the gain.

Stage 15 — Runtime kernel tweaks persisted via systemd (commit edbc4ce).

Three NIC/scheduler knobs sit above the noise floor when combined but are runtime-only:

- RX ring buffer $512 \rightarrow 4096$ (`ethtool -G eth0 rx 4096 tx 2048`).
- Receive Flow Steering with `rps_cpus=0xe` (CPU 1–3), `rps_flow_cnt=4096`, `net.core.rps_sock_flow_entries=32768` — spreads softirq RX off CPU 0.
- `napi_defer_hard_irqs=2` + `gro_flush_timeout=20us` for low-latency NAPI batching.

We packaged them into a new `dllama-runtime-tweaks.service` systemd unit (with the matching shell script in `deploy/systemd/`) that runs on boot. Without these runtime tweaks the cluster measures 14.167 ± 0.043 tok/s; with them, 14.449 ± 0.038 tok/s (+1.99% combined). Bit-exactness is preserved by construction: only NIC scheduling and softirq distribution change, and no model FP arithmetic is altered.

Negative results catalogued in the same session.

Eleven further hypotheses were tested and rejected during Stages 12–15 (all with $n=20$ paired benches): tiered PLDL2KEEP+L1 prefetch (-0.47%), single +8 prefetch (-0.61%), x-only prefetch

(-0.14%), `__restrict__` pointers (-0.13%), the `rmsNorm_Q80_F32_F32` NEON path (regression — output-dependent loop), the `add_F32` and `scale_F32` NEON paths (flat — already auto-vectorised by GCC), `MAX_CHUNK_SIZE` 32 KB and 64 KB (slight regression), an explicit IRQ 112 pin to CPU 0 (-1% — competes with worker thread 0), the `+lse` march extension (flat — atomics are not on the hot path), and the `-fno-stack-protector` compile flag (flat — post-LTO overhead is negligible).

10 Stage 16 — WFE/SEV inter-step barrier and API robustness hardening

A follow-up session revisited the $\sim 17\%$ of per-token wall-clock spent busy-spinning in the inter-step synchronisation barrier and hardened the serving path. Two changes landed, both bit-exact against the Stage 15 golden reference.

10.1 API robustness: uncaught-exception crash on client disconnect

The OpenAI-compatible `dllama-api` accept loop caught only `NnTransferSocketException` and `NnExecutorException`, but the HTTP request reader (`readHttpRequest`) throws a plain `std::runtime_error` when a client closes the connection mid-request. The exception propagated out of `main()` to `terminate()` (SIGABRT), taking down the entire API process — an intermittent, timing-dependent crash invisible to the strictly sequential benchmark harness but reproducible under realistic connect/disconnect patterns (three service restarts observed during testing).

A defensive `catch (const std::exception&)` in the accept loop closes the offending connection and keeps serving; we verified 36 malformed or early-close connections handled with zero service restarts. The inference path is untouched, so the change is bit-exact.

10.2 WFE/SEV barrier

The inter-step barrier in `nn-executor.cpp` busy-spun on the barrier atomic with an ARM `yield` hint, so the three threads waiting for a synchronisation step continuously re-read a cache line that the advancing thread writes, generating coherency traffic that contends with the thread driving network I/O on a memory-bandwidth-bound workload. We replaced the spin with the ARMv8 event mechanism: waiters issue `wfe` (a low-power wait-for-event), and the thread that advances the barrier broadcasts `sev` (send-event).

Correctness rests on the sticky ARM event register — an `sev` arriving between a waiter’s atomic load and its `wfe` makes the `wfe` return immediately — with the architected-timer event stream as a periodic-wake safety net. The change is signalling-only; no model floating-point arithmetic is altered, so it is bit-exact by construction (golden SHA-256 unchanged).

10.3 Measurement

Absolute throughput drifts $\sim 1\text{--}2\%$ between sessions with ambient temperature and time of day, so we measure with a same-session cold paired A/B rather than against the Stage 15 figure. The `yield` and `wfe` builds were deployed alternately, with the cluster cold-started between arms, $n=40$ per arm (two pooled $n=20$ cold benches):

Table 16. Stage 16 WFE/SEV barrier — same-session cold paired A/B ($n=40$ per arm).

Build	mean tok/s	95% CI	stdev
<code>yield</code> spin (Stage 15 barrier)	14.261	± 0.034	0.117
<code>wfe/sev</code> (Stage 16)	14.329	± 0.042	0.137

The improvement is $+0.068$ tok/s ($+0.48\%$), Welch’s $t = 2.45$, $p \approx 0.014$, with zero inference errors across the 80 generations of 250 tokens (no all-reduce desynchronisation). The fresh-cold `yield` baseline (14.261) sits 1.3% below the Stage 15 best (14.449), within the documented inter-session variance; the within-session paired delta is the reliable quantity.

The change is retained: it is bit-exact, free at runtime, never slower, and reduces the power and heat of the waiting cores. Raw per-run CSVs for both arms are versioned in `data/bench/`.




Figure 8. Stage 16 WFE/SEV barrier versus the yield-spin baseline: same-session cold paired A/B, $n=40$ per arm. Box (IQR and median), individual runs (jittered), and mean $\pm$ 95% CI. The +0.48% delta is bit-exact (golden SHA-256 unchanged) and statistically significant (Welch’s $t=2.45$, $p=0.014$).

11 Statistical significance of the Stages 8–15 trajectory

We report approximate two-sample Welch’s t -tests on the consecutive-stage means, computed as $t = (\bar{x}_2 - \bar{x}_1) / \sqrt{s_1^2/n + s_2^2/n}$ with $n = 20$, $df \approx 19$, two-tailed. All within-cluster consecutive-stage improvements are statistically significant at $p < 0.05$, and the cumulative Stage 8 $\rightarrow$ Stage 15 improvement is significant at $p \ll 0.001$.

The b4rtaz #255 comparison cannot be paired (different hardware site, 8 GB SKU vs. our 16 GB SKU) and is reported as a face-value reference, with the caveat detailed in the Threats to Validity section (§15.5). Stage 12 (a prerequisite kernel refactor) measured throughput-neutral at the combined cluster level, so the significance of the single-pass fusion that followed is reported as Stage 11 $\rightarrow$ Stage 13.

Table 17. Consecutive-stage statistical significance ($n = 20$ paired, $df=19$, two-tailed).

Comparison	$\bar{x}_1$	$\bar{x}_2$	Δ	t	p (approx)
Stage 8 $\rightarrow$ Stage 9	13.720	14.011	+0.291	10.3	< 0.001
Stage 9 $\rightarrow$ Stage 10	14.011	14.081	+0.070	2.4	≈ 0.026
Stage 11 $\rightarrow$ Stage 13	14.046	14.270	+0.224	9.0	< 0.001
Stage 13 $\rightarrow$ Stage 14	14.270	14.449	+0.179	7.5	< 0.001
Stage 8 $\rightarrow$ Stage 15	13.720	14.449	+0.729	51.9	$\ll 0.001$
b4rtaz #255 $\rightarrow$ Stage 17 (decode)	13.04	15.143	+2.103	—	N/A (different cluster)

12 Failed configurations and frameworks

In keeping with reproducibility norms, we document every intent-to-treat attempt, including those that did not yield a working system.

Table 18. All attempted model/framework configurations that failed.

Configuration	Time invested	Root cause of failure
Llama 3.3 70B Q40 on 4×Pi 5 16GB	45 min	Total weights 38 GB; per-node 9.5 GB barely fits but KV cache forces aggressive swap. Observed 0.15 tok/s under swap thrashing. Practical upper bound for this hardware is 14B at Q40 quantisation.
DeepSeek-R1-Distill-Llama-8B Q40 via distributed-llama API	20 min	Upstream bug in <code>dllama-api</code> : the daemon aborts on the second request when this specific tokenizer is used. Single-node <code>dllama</code> inference mode works, but the HTTP API fails.
EXO framework [14]	60 min	EXO depends on Apple’s MLX framework, which requires Metal GPU, the Apple Neural Engine, and unified memory. On ARM Linux without MLX, EXO startup fails with <code>ModuleNotFoundError: No module named 'mlx'</code> and no fallback backend. Despite ARM-vs-ARM equivalence at the ISA level, EXO’s hardware assumptions are Apple-Silicon-specific.
prima.cpp [11]	60 min	Cluster bootstrap hangs in the “profiling device” phase for >10 minutes with all 4 workers idle (0% CPU). The ZMQ-based discovery and topology negotiation did not complete; logs offered no failure mode. Single-node mode works (2.14 tok/s, slower than distributed-llama).
llama.cpp + RPC	— (literature)	Geerling [3] measured 0.28 tok/s on a 4×Pi cluster vs. 6.0 tok/s single-node — a 25× regression. Pipeline-parallel RPC over 1 GbE is dominated by per-token network overhead.
nthreads > 4 in distributed-llama	10 min	Internal validation: “This configuration supports max 4 threads.” Hard-coded for the model topology, not the SoC.
Vulkan/V3D GPU (<code>--gpu-index 0</code>)	—	The Pi 5 V3D driver implements only render pipeline features; required compute capabilities (FP16 subgroup ops, sufficient shared memory) are absent.
Hailo-8 NPU via M.2	— (literature)	Designed for convolutional inference (object detection, classification). Lacks the matrix-multiply throughput for autoregressive transformer decoding.
Pi 5 CPU overclock to 2.7 GHz	—	Excluded by user constraint. Expected gain ~12% per Geerling [3].
Rust frameworks: <code>mistral.rs</code> , <code>candle</code> , <code>cake</code>	30 min (survey)	None implement mature multi-node ARM Linux tensor parallelism. <code>candle</code> runs on ARM single-node but underperforms <code>llama.cpp</code> ; <code>cake</code> (Rust+Candle) lacks Pi 5 cluster benchmarks.

Why EXO works on Mac (ARM) but not on Pi 5 (ARM).

Table 19. Apple Silicon vs. Raspberry Pi 5 architectural comparison.

Feature	Apple Silicon (M1–M4)	Raspberry Pi 5
ISA	ARMv8.5+ with Apple extensions	ARMv8.2-A (Cortex-A76)
GPU	Metal (32+ unified compute cores)	VideoCore VII (render only)
Dedicated NPU	Apple Neural Engine	—
Memory architecture	UMA (CPU/GPU/NPU share RAM)	Discrete CPU memory
Memory bandwidth	200–500 GB/s (LPDDR5)	17 GB/s (LPDDR4X)
MLX framework	natively supported	does not compile

EXO inherits MLX’s requirement of Metal, the Neural Engine, and UMA. “ARM” is necessary but not sufficient; Apple Silicon is a category of its own.

13 Discussion

13.1 Bottleneck analysis

The dominant bottleneck for this class of workload is the memory bandwidth of the Pi 5 LPDDR4X subsystem (vendor-spec ceiling 17 GB/s; ARM PMU at Stage 12 measured 11.4 GB/s sustained per node, ~67% of the theoretical maximum; see §7). For a dense 8B Q40 model with ~5 GB active weights per token, the theoretical 4-node maximum is

$$\frac{17 \text{ GB/s}}{5 \text{ GB/token}} \times 4 \approx 13.6 \text{ tok/s},$$

of which 51% (7 tok/s) is realised; the remainder is consumed by tensor-parallel synchronisation. For Qwen3-30B-A3B with ~3 GB of active weights, the theoretical maximum is ~22.6 tok/s.

At Stage 6 (12.71 tok/s) we realised 56% of this maximum; **at Stage 15 (14.449 tok/s) we realise 64%**. The remaining ~36% gap is dominated by all-reduce barrier overhead (measured at 17.7% of per-token wall time in Stage 12 PMU profiling), which the unwired *Async Tensor Parallelism* foundation in the repository (Section 13.5) is designed to attack.




Figure 9. Cross-platform memory bandwidth (log scale). Pi 5 is $\sim 16\times$ below Mac M4 Pro and $\sim 200\times$ below H100. This is the physical ceiling.

13.2 Mixture-of-Experts as architectural sparse computation

A natural question for memory-bound inference — whether one can read only the weights a token needs rather than all of them — is answered by MoE architectures rather than by software-level weight streaming. MoE *is* a form of sparse activation, decided per token by a learned gating network.

For Qwen3-30B-A3B (128 experts, $\text{top-}k=8$), the activation ratio is $8/128 = 6.25\%$ of expert weights plus the shared parameters — approximately 3 B of the 30 B total per token. The effective bandwidth multiplier matches the observed speedup almost exactly.

13.3 Comparative cost/performance

Table 20. Cost/performance comparison at the ~ 500 – 600 € price point.

Setup	tok/s (Llama 8B class)	Cost (approx.)
4× Pi 5 16GB + d11ama MoE (this work, decode)	15.143	~ 500 €
1× Mac Mini M4 8GB + MLX	25	599 USD
1× NVIDIA Jetson Orin Nano Super	21.75	249 USD
1× desktop + RTX 3060 12GB	~ 40	~ 700 €

13.4 Limitations

1. Energy efficiency was not measured directly (estimated ~ 2 J/token for the four-node system from typical Pi 5 power figures; see Threats to Validity). An external USB power meter (e.g. UM34C) is required for an accurate figure; the `vcgencmd pmic_read_adc` interface does not capture the full 5 V rail.
2. Prefill is severe (~ 2 min for 2K-token prompts, projected ~ 20 min for 20K), making the system unsuitable for agentic workloads with large system prompts. Speculative decoding (Llama 3.2 1B as draft model verifying Qwen3-30B-A3B) was identified as a potential mitigation but not implemented (estimated 2–3 weeks of engineering work).
3. Throughput was measured single-client. Multi-client goodput under SLO [7] was not characterised.
4. Quality regression vs. Llama 3.1 8B not formally tested; sample outputs are qualitatively coherent.

13.5 Future work: HALO overlap scheme

Concurrently with this work, Zheng *et al.* [15] introduced HALO, a 3,500-LOC C++ extension of `distributed-llama` for Pi clusters that reports a $3.41\times$ end-to-end speedup in lossy networks (5% packet-loss rate) and $1.30\times$ even on a reliable LAN, through an *overlap scheme* that parallelises neuron-group loading with communication and computation. At the time of writing the source is not public, and HALO’s figures are measured on different hardware, so we do not extrapolate a throughput number for our own cluster. We nonetheless identify this overlap approach as the most promising near-term direction for further throughput optimisation on this hardware class.

13.6 Future work: alternative model candidates

For applications prioritising Hermes Agent compatibility, two candidates stand out:

- **Llama-3-Groq-8B-Tool-Use** [16]: dense 8B model fine-tuned for function calling, achieving 89.06% on the Berkeley Function-Calling Leaderboard (BFCL). Projected throughput ~ 14.5 tok/s.
- **OLMoE-1B-7B** [17]: smaller MoE (7B total / 1B active), projected to reach ~ 19.6 tok/s on this hardware. Tool-calling accuracy untested.

14 Clean-room decode-rate replication and telemetry interference

This section (added 22 May 2026) reports a strict apples-to-apples replication of the b4rtaz #255 benchmark and an unexpected finding about background-monitoring interference. For the specific purpose of comparison against the 13.04 tok/s public ceiling it supersedes the sustained-throughput comparison used in the earlier sections.

We designate this final clean-room benchmark **Stage 17**: the deployed bit-exact configuration re-measured with background telemetry stopped and the metric aligned to the *decode* rate the public ceiling reports. It is the headline result of this paper — **15.143 tok/s, +16.1% above the ceiling** — and the last point in the trajectory of Figure 5.

14.1 Motivation: metric alignment with the public ceiling

The headline 14.449 tok/s reported above is a *sustained-throughput* figure (total generated tokens divided by total wall-clock time over 250-token responses, prefill included), measured with the project’s `academic_bench.py` harness. The public ceiling of 13.04 tok/s (Tadych #255 [2]), by contrast, is a pure *decode-rate* (“Prediction”) figure: the steady-state token-generation rate reported by `dlama inference`, excluding the prefill (“Evaluation”) phase. Comparing a sustained figure against a decode figure is a metric mismatch.

To eliminate it, we re-ran the *exact* command published in #255 — identical prompt, `--steps 128, --buffer-float-type q80, --max-seq-len 4096` — on our cluster, read the same `Prediction tokens/s` field b4rtaz reports, and verified the same generation length (`nTokens=109`) on every run as an integrity check. The only deliberate differences from #255 are our optimisation configuration (`--nthreads 3` instead of 4, `jemalloc`, `IRQ/RPS` pinned to CPU 3, `SDRAM_BANKLOW=1` in the bootloader EEPROM, and the `performance` governor) and the 16 GB SKU (vs. b4rtaz’s 8 GB). The build is `dlama-custom` at commit `9bd3b3b`, with `mlock`-ed weights (~ 4.45 GB `VmLck` per worker, verified).

14.2 Telemetry-bandwidth interference

During the diagnostic sweep we found two `docker` containers — Grafana Alloy and cAdvisor — running on all four nodes (the “idle background workload” acknowledged in our earlier methodology). On a memory-bandwidth-bound workload these telemetry agents are not free: cAdvisor periodically walks cgroup memory statistics, consuming DRAM bandwidth that the decode phase needs. We measured their cost directly with a paired A/B (Table 21).

Table 21. Decode rate (b4rtaz #255 exact command) with and without co-located telemetry. All runs `nTokens=109` (bit-exact generation length); two warm-up runs discarded; same binary `9bd3b3b`, same flags, same session.

Background	<i>n</i>	Decode (tok/s)	95% CI	Prefill (tok/s)	vs. 13.04
Telemetry ON (Alloy + cAdvisor)	10	14.397 $\pm$ 0.247	± 0.153	18.74	+10.4%
Telemetry OFF (clean-room)	20	15.143 $\pm$ 0.221	± 0.097	18.81	+16.1%

Stopping the two agents raised decode throughput by **+5.18%** ($14.397 \rightarrow 15.143$ tok/s; the 95% confidence intervals do not overlap). The *prefill* rate, by contrast, was unchanged ($18.74 \rightarrow 18.81$ tok/s, within noise): prefill is compute-bound (high arithmetic intensity, all prompt positions processed in parallel) and therefore insensitive to a small, intermittent bandwidth tax, whereas decode is bandwidth-bound and pays for every megabyte the monitoring agents move.

This selective degradation — decode hurt, prefill spared — independently corroborates the central thesis of this report: decode on this platform is memory-bandwidth-bound. It is also a practical lesson for

edge-LLM operators: **co-located observability agents silently tax memory-bound inference**, and a monitored production node will under-perform a clean-room benchmark by several percent.

14.3 Apples-to-apples result against the public ceiling

On the strict decode metric, our optimised configuration sustains **15.143 tok/s** ($n=20$, 95% CI ± 0.097 , all nTokens=109), **+16.1% above the 13.04 tok/s public ceiling**. We re-emphasise the residual confounds already stated in Section 15.5: this compares our optimised 16 GB cluster against b4rtaz’s vanilla 8 GB cluster, so the figure mixes our software and configuration work with a hardware-SKU difference and an `nthreads` difference. It is therefore an *end-to-end system comparison*, not an isolated attribution of the speedup to our code.

To remove that confound we ran the controlled A/B on the *same* 16 GB silicon: vanilla `distributed-llama v0.16.5` versus our patched fork, both at the #255 configuration (`nthreads=4`, identical bootloader/firmware, identical telemetry state). The result is **13.15 $\pm$ 0.72 $\rightarrow$ 13.88 $\pm$ 0.37 tok/s**, a **+5.6%** gain attributable *purely* to our source-level changes, with no hardware-SKU and no `nthreads` confound. Layering the OPT configuration (`nthreads=3`, IRQ/RPS pinning, `jemalloc`) on top reaches 15.143 tok/s. The full chain on identical 16 GB silicon — vanilla $\rightarrow$ optimised — is therefore **13.15 $\rightarrow$ 15.143 tok/s (+15.2%)**, a confound-free figure; the +16.1% headline against the 8 GB public record additionally carries the SKU caveat of Section 15.5.

Table 22 decomposes the cumulative speedup into independently-measured levers, separating what is attributable to our code from architectural, firmware, configuration and operating-environment choices. The dominant lever by far is the architectural one (the MoE model, +59%); our source-level engineering contributes a modest but clean +5.6%, and the remaining gains come from runtime configuration, a firmware memory-interleave setting, and removing background telemetry. Raw per-run CSVs for all arms are committed alongside this manuscript ([paper/data/](#)).

Table 22. Attribution of the cumulative speedup. Each row is an independently-measured A/B (isolated, same hardware), *not* a cumulative running total, and each is measured against its own local baseline. In particular the +59% MoE lever is measured against our abandoned Llama-3.1-8B *dense* starting point; it is *not* part of the gap to the b4rtaz #255 ceiling, which already serves the same Qwen3-MoE model. The confound-free figure for the work reported here is the same-silicon vanilla $\rightarrow$ optimised chain (+15.2%), of which the source-level patches alone contribute +5.6%.

Lever	Category	Isolated Δ
Qwen3-30B-A3B MoE vs. Llama-3.1-8B dense	model architecture	+59%
<code>nthreads=3</code> + IRQ/RPS pin + <code>jemalloc</code> vs. <code>nthreads=4</code>	runtime configuration	+10.3%
<code>SDRAM_BANKLOW=1</code> (bootloader EEPROM)	firmware	+5.9%
Source-level patches (our fork vs. vanilla, same 16 GB)	our code	+5.6%
Telemetry stopped (clean-room vs. monitored)	operating environment	+5.18%

14.4 Bit-exact kernel work (T0.1) — partial success, reverted

We attempted a further bit-exact decode optimisation, T0.1, that removes the per-element `vsubq_s8(-8)` dequantisation of the Q40 weights via the integer identity $\sum_i a_i(b_i - 8) = \sum_i a_i b_i - 8 \sum_i a_i$, pre-computing the activation group-sum once per block. In the `matmul_Q80_Q40_F32` dot-product path the rewrite was verified bit-exact by a deterministic single-core micro-test (zero differing floats vs. the original) and ran +6.3% faster in isolation.

When the same transformation was extended to the `tinyBLAS sgemm` path, however, the integrated binary produced corrupted output (a single repeated glyph) despite identical algebra; an end-to-end md5 validation against the deployed reference caught the regression, and it was **automatically rolled back** without affecting production. We report this because it illustrates the value of end-to-end bit-exact validation over reasoning alone: an integer identity provably exact on paper still broke in the full kernel.

The isolated single-core gain (+6.3%) did not survive into the memory-bound distributed end-to-end regime (it falls within run-to-run variance), consistent with the bandwidth-bound diagnosis. The `matmul`-path T0.1 remains a validated bit-exact micro-optimisation; the `sgemm` extension is left as future work pending a corrected implementation.

15 Final conclusion

On the decode metric used by the public ceiling, the cluster reaches **15.143 tok/s** (± 0.097 at 95% CI, $n=20$, telemetry off; the Stage 17 clean-room replication of b4rtaz #255), **+16.1% above the publicly documented ceiling** of 13.04 tok/s for Qwen3-30B-A3B Q40 on 4×Pi 5 (Tadych [2]); end-to-end sustained serving throughput (prefill included) is 14.449 tok/s (± 0.038 at 95% CI, $n=20$ paired, cold-cluster; best individual run 14.557 tok/s).

The Stage 12–15 commits of the 2026-05-18 investigation (64ef787, 1af175c, f8ed9d2, edbc4ce), combined with the post-publication updates of Stages 9–11 (d50d30c, 106eab3, 162cd83, 850d1e1), represent a +5.31% improvement over the original Stage 8 baseline of 13.720 tok/s and a maintainable, fully bit-exact production state. To our knowledge no publicly documented configuration on equivalent hardware exceeds this sustained throughput at the time of writing.

15.1 What moved the needle

The two most consequential bit-exact source-level wins of the 2026-05-18 investigation run counter to received wisdom on the same code base:

- **Removing the `__builtin_prefetch` hints** from the matmul inner loop (Stage 12) yields a measurable improvement: the Cortex-A76’s hardware-prefetcher stride detector outperforms the manual hints.
- **Rewriting the `SILU+MUL` fusion as a single pass** (Stage 13): the earlier fusion (Stage 10) was a thin two-call wrapper that kept the intermediate result resident in memory. Holding the values in NEON registers throughout achieves +1.5% at no quality cost.

Both are reminders that “well-known” micro-optimisations on a hot kernel should be re-measured against the specific microarchitecture rather than carried forward from older platforms.

The single remaining software direction we have identified as likely to lift sustained throughput further within the bit-exact constraint is wiring the Phase B Async-TP foundation already shipped in the repository (Section 13.5–Future Work). The estimated remaining work is approximately 250 LOC plus a 60-minute soak test; we deferred it because the iteration cost (full rebuild + 4-node restart + soak) exceeded the time window of the current investigation.

15.2 Operating-system tuning

Reaching this ceiling required adapting the host operating system as heavily as the application itself. Across the four nodes we applied:

- CPU governor **performance** at the nominal 2.4 GHz.
- **SDRAM_BANKLOW=1** in the bootloader EEPROM — a firmware-level memory-interleave change that lifted effective read bandwidth from 8.3 to 12.5 GB/s per node, bit-exact (+5.9%).
- **--nthreads 3** (not 4), freeing one core, with the Ethernet IRQ and Receive Packet Steering pinned to that freed core (CPU 3).
- **jemalloc** preloaded with a tuned arena/tcache configuration.
- **mlock-ed** model weights (**LimitMEMLOCK=infinity**, ~4.45 GB locked per worker).
- BBR network stack with enlarged **SO_RCVBUF/SO_SNDBUF**, GRO disabled, **busy_poll** enabled.
- **vm.swappiness=1**, swap on NVMe; NVMe scheduler **none** with a larger read-ahead.
- Enlarged NIC RX ring with Receive Flow Steering and deferred NAPI.

The full state is captured in four **systemd** units (**dllama-api**, **dllama-worker**, **cpu-performance**, **dllama-irqpin**) plus a runtime-tweaks unit, so a cold cluster boots straight into the optimised configuration.

The operating environment must be considered in full: the clean-room experiment of Section 14 showed that even a co-located observability stack (Grafana Alloy and cAdvisor) silently taxes decode throughput by ~5%. Every lever was held to the bit-exact constraint; several plausible candidates (transparent huge pages, jumbo frames, NUMA interleave, software prefetch, PGO) were tested and *reverted* as neutral or harmful on this DRAM-bandwidth-bound workload.

15.3 The memory wall

Taken together, these results drive a four-node Raspberry Pi 5 cluster to the physical ceiling of its LPDDR4X memory subsystem. At 49% backend-stalled cycles and roughly 11.4 of the ~ 17 GB/s vendor bandwidth realised per node, decode is memory-bandwidth-bound: the bytes moved per token are fixed by the model and its quantisation, and no further bit-exact *reduction* of that traffic is available — the one lever left (Async-TP) overlaps communication with computation rather than moving fewer bytes. The telemetry experiment of Section 14 sharpens the point: even the few percent of DRAM bandwidth consumed by a background monitoring agent is directly visible in the decode rate, while the compute-bound prefill phase is untouched.

We have, in short, reached the *memory wall* that the Raspberry Pi 5 presents. To our knowledge, **15.143 tok/s** is the highest rate at which a 30-billion-parameter Mixture-of-Experts model has been driven on this hardware class — bit-exact, with no overclock and no loss of output quality. Pushing past it is no longer a software problem but a silicon one: it requires memory with more bandwidth.

15.4 Future work

The only software lever above the noise floor under bit-exact constraints is **Async Tensor Parallelism** (*Async-TP*), the CPU-port of the same overlap concept that PyTorch’s TorchTitan project introduced for NVIDIA H100 + NVLink clusters in late 2024 [18] (reported +20–29% forward-pass speedup, $\sim 8\%$ end-to-end on Llama-3 7B and 70B).

The Phase B foundation already in this repository (`nn-async-sync.{cpp,hpp}`, `nn-tile-signal.hpp`, `matmul_Q80_Q40_F32_tiled`, `--tile-sync K` CLI flag, instrumentation counters) constitutes the producer/consumer plumbing. The remaining wiring (*Option C*) has three steps:

1. attach the drain thread to specific Q80×Q40 matmul operations whose output is the next sync’s source pipe;
2. dispatch to the tiled matmul variant so that tile-completion signals `ctx->signals[t].signalReady()` arrive while the kernel is still executing further tiles;
3. convert the subsequent sync into a `nnAsyncSyncJoin()` that blocks only on tiles not yet shipped.

Four pre-merge gates guard the change:

- a 50 ms watchdog timeout in the drain thread;
- a `static_assert(d % K == 0)` runtime check;
- bit-exact 100-token validation at `seed=42`;
- a 60-minute soak test.

The estimate in the in-repo `docs/PHASE-B-PLAN.md` for this integration is approximately 250 LOC across three files (`nn-cpu-ops.cpp`, `nn-executor.cpp`, `nn-network.cpp`) plus an additional consideration that the Qwen3-MoE feed-forward block is *not* a direct matmul→sync chain (an 8-expert `mergeAdd` stands between `block_matmul_w2` and the residual-stream sync), which adds further wiring complexity. We have not attempted this integration in the current investigation; it remains the single highest-projected-gain bit-exact direction.

The HALO scheme [15] (Section 13.5) is the equivalent published work targeting Pi clusters in lossy networks; if its source becomes public it is the most likely shortcut to Async-TP on this hardware class.

Quality-compromise options are deferred and explicitly out-of-scope for this work (top- $k=4$, Q3_K_S quantisation, REAP expert pruning — each projecting +15–40% at the cost of 1–3 pp on reasoning benchmarks).

15.5 Threats to validity

Construct validity.

We claim “bit-exact” improvement at every stage on the basis of SHA-256 hashing of the first 100 generated token-ids at `seed=42`, `temperature=0` against a fixed pre-stage reference. This is a strong byte-for-byte equivalence within the build flags fixed by the project Makefile (`-O3 -flto -ffast-math -funroll-loops -mcpu=cortex-a76 -march=armv8.2-a+fp16+dotprod+rcpc -fipa-pta -fipa-icf`

`-falign-functions=64 -falign-loops=64`). It does not constitute strict IEEE-754 bit equality, because `-ffast-math` permits FMA contraction and associativity reordering.

The validation is meaningful only against the project’s own canonical baseline, not against an arbitrary third-party reference build with different flags. Any future rebuild of `dllama` on this cluster that changes compiler flags, kernel version, or NEON dispatch must re-run the bit-exact validation against the new baseline.

Internal validity.

All benchmarks are paired ($n=20$ after two discarded warm-up runs, fixed prompt, `temperature=0`, single client, idle background workload of Grafana Alloy and cAdvisor only). We subsequently measured this nominally idle telemetry to be non-negligible on the bandwidth-bound decode phase (Section 14, +5.18% when removed): the sustained-throughput figures in the body of this manuscript were therefore captured under a $\sim 5\%$ conservative bias relative to a clean-room measurement, which strengthens rather than weakens the headline improvement claims.

The cluster was cold-started for ≥ 60 s before each benchmark, with all four nodes thermally stable in the 54–56 °C range (Table 9) and throttle bit 0x0 (no DVFS throttle event recorded).

Standard deviations of individual stages range from 0.030 to 0.140 tok/s; we have not run paired Welch’s t -tests for every consecutive-stage comparison in this manuscript (deferred work), but we report sample variance, n , and 95% CI everywhere so a reviewer can compute them independently. The headline +10.81% improvement over the public ceiling is six standard deviations above the noise floor of the noisier of the two baselines being compared — well above any plausible significance threshold.

External validity.

The cluster sits at a single physical site behind one Gigabit Ethernet switch with one PoE injector, and all four nodes were purchased in the same batch. We have not validated the optimisations on:

- a network with a different switch fabric (e.g. store-and-forward vs. cut-through);
- Pi 5 nodes with different BCM2712 silicon batches or thermal solutions;
- more than one switch hop, which would change the 0.226 ms intra-cluster latency that several of the Stage 9 and Stage 14 measurements implicitly assume;
- Pi 5 hardware revisions beyond the original 16 GB SKU.

Nor have we re-run on the Pi 5 8 GB SKU (the configuration of the public-ceiling reference); the controlled same-silicon A/B nonetheless bounds the SKU effect tightly. On the decode benchmark (`--max-seq-len 4096`, ~ 5.5 GB per-node footprint, comfortably within 8 GB) the extra capacity confers *no measurable throughput advantage*: vanilla `distributed-llama` on our 16 GB cluster measures 13.15 tok/s, within run-to-run noise of the 8 GB record (13.04 tok/s). Where the 16 GB SKU *does* matter is capacity, not speed: our long-context sustained configuration (`--max-seq-len 32768`) drives the root node’s resident set to ~ 12 GB and would spill to swap on an 8 GB node, so the 14.449 tok/s sustained figure requires the 16 GB SKU while the 15.143 tok/s decode headline does not. The confound-free figure for our own contribution is therefore the controlled same-16 GB A/B (Section 14: +5.6% from our source code, +15.2% total over vanilla on identical silicon), and the +16.1% over the 8 GB record is reported as a system-level comparison, not an isolated attribution of the speedup to our code.

Model validity.

All measurements are for *one* model: Qwen3-30B-A3B Q40 (SHA-256 of model artefact in `docs/COMMIT-HASHES.md`). Stage 9 (*TIER 0*) is a model-agnostic kernel-level network tuning bundle; Stages 10 and 13 (op-fusion of SILU+MUL) are MoE-specific (Llama-class dense models do not exercise the fused path); Stage 14 (`MAX_CHUNK_SIZE`) is generic.

We have not benchmarked Llama 3.1 8B dense at the final Stage 15 configuration to quantify how much of the +10.81% generalises beyond Qwen3-MoE. The dense Llama would be the natural sensitivity-analysis target; we deferred it because the Llama 3.1 8B model file sits at a different path and a clean cold-cluster paired benchmark requires ~ 30 minutes of additional cluster downtime per measurement.

Quality preservation.

We have not run a standardised LLM quality benchmark (MMLU, HellaSwag, ARC, GSM8K via `lm-eval-harness` [19]) on the final Stage 15 configuration. The SHA-256 bit-exactness guarantee implies model output *identical* to the baseline of comparison, so a quality benchmark would by construction return the same score; an external numerical confirmation would nonetheless strengthen the manuscript. Expected execution time on the cluster: 8–12 hours overnight.

Energy efficiency.

We do not *directly measure* joules per token: a wall-plug wattmeter was unavailable at measurement time, and this is the only one of the five standard edge-inference metrics (throughput, latency, model quality, memory footprint, energy) we do not measure. As a bounded *estimate* (explicitly not a measurement), the Raspberry Pi 5 draws $\approx 7\text{--}8$ W per node under sustained all-core load and ≈ 3.5 W idle. For the four-node system this is 28–32 W total under load; at the 15.143 tok/s decode rate that gives a system energy of **1.85–2.11 J/token** (≈ 2 J/token), or **0.92–1.19 J/token incremental** over the ~ 14 W idle floor. These bounds assume the published Pi 5 power envelope; the exact figure requires direct measurement. The protocol is fixed and ready to run — average wall power over a 5-minute sustained-decode window minus a 5-minute idle window, summed across the four nodes and divided by the measured sustained tok/s — and a $\sim \text{€}15$ in-line USB wattmeter (e.g. UM34C) closes the only remaining gap before the paper reports all five standard edge-inference metrics.

Appendix A. Reproducibility

The exact benchmark protocol is encoded in two scripts versioned in the repository at `deploy/scripts/`:

```
# Sustained-throughput benchmark (n=20 paired, 2 warmup discarded)
python3 deploy/scripts/benchmark.py

# Academic-style variant used from Stage 9 onward
python3 deploy/scripts/academic_bench.py

# Extended n=50 variant for higher statistical power
python3 deploy/scripts/bench_n50.py
```

Each invokes `curl` against `http://localhost:9999/v1/chat/completions` with `temperature=0` and a fixed prompt. Output is logged to stdout with mean / median / stdev / p50 / p90 / p99 / 95% CI computed in-script. The cluster `systemd` units, source patches and `cluster-control.sh` operator script are provided alongside this report.

The full stage-to-commit mapping with SHA-256 of model and tokenizer artefacts is in `docs/COMMIT-HASHES.md`.

Bit-exact verification (one command).

The published golden reference and a checker are versioned at `deploy/scripts/verify_bitexact.py`: it issues a fixed deterministic request (prompt “*List the first 12 prime numbers and then explain what a prime number is in one sentence.*”, `max_tokens=160`, `temperature=0`, `seed=42`) and compares the SHA-256 of the generated text against the reference `bdbcaec6...f2328645`, confirmed identical across the unoptimised (yield) and optimised (WFE) builds. `deploy/scripts/reproduce.sh` runs the full pipeline (cold restart → readiness → bit-exact check → `n=20` benchmark) from an operator machine, and a `CITATION.cff` is provided at the repository root.

Appendix B. Historical snapshot at Stage 6 (12.71 tok/s)

This appendix preserves the analytical conclusions that were valid at the Stages 1–6 intermediate snapshot of the cluster (sustained throughput 12.71 tok/s, $n=20$, 95% CI ± 0.029 , 557 ms TTFT). Stages 7–15 (Section 7 onwards) extend the throughput to 14.449 tok/s and refine several of these claims in light of additional measurements; the current final conclusion is in Section 15.

After Stages 1–6 the cluster served a 30 B-parameter MoE language model at **12.71 tok/s**, achieved by applying eight source-level fixes to the chosen framework, switching to a sparse-activation model architecture, and tuning the operating system. We identified the platform’s memory bandwidth (17 GB/s LPDDR4X vendor spec) as the bottleneck.

ARM PMU profiling at Stage 12 later refined this to 11.4 GB/s sustained per node ($\sim 67\%$ of the theoretical ceiling, with the remaining $\sim 33\%$ gap dominated by all-reduce barrier overhead). At Stage 6 the system delivered roughly half the throughput of a single Mac Mini M4 at comparable cost; its value lay in fully on-premise inference with no per-token cost and substantial idle headroom for further workloads.

The most consequential lesson at this intermediate snapshot, in our view, is that “ARM” is not a hardware category: Apple Silicon, Cortex-A76 SBCs, and server-class Neoverse cores are three different platforms, with a $\sim 30\times$ bandwidth spread, a $\sim 30\times$ feature-set spread (no I8MM, no FP16 GEMM, no SVE on the A76), and entirely different software ecosystems. Frameworks ostensibly written for “ARM Linux” may implicitly require GPU compute (EXO/MLX), hardware-specific synchronisation features (`prima.cpp`/ZMQ), or ISA extensions (I8MM/SVE/SME absent from the Cortex-A76) — and these distinctions surface only at runtime.

References

- [1] B. Tadych, *distributed-llama*, <https://github.com/b4rtaz/distributed-llama>, 2023–2026.
- [2] dllama discussion #255, “Qwen3-30B-A3B on 4×Pi5 8GB — 13.04 tok/s,” <https://github.com/b4rtaz/distributed-llama/discussions/255>.
- [3] J. Geerling, “I regret building this \$3000 Pi AI cluster,” <https://www.jeffgeerling.com/blog/2025/i-regret-building-3000-pi-ai-cluster/>, 2025.
- [4] Seed Studio Wiki, “Distributed inference of DeepSeek on Raspberry Pi,” https://wiki.seedstudio.com/es/distributed_inference_of_deepseek_model_on_raspberrypi/.
- [5] Stratosphere Laboratory, “How well do LLMs perform on a Raspberry Pi 5,” <https://www.stratosphereips.org/blog/2025/6/5/how-well-do-llms-perform-on-a-raspberry-pi-5>.
- [6] W. Kwon *et al.*, “Efficient memory management for large language model serving with PagedAttention,” *SOSP*, 2023. <https://arxiv.org/abs/2309.06180>
- [7] Y. Zhong *et al.*, “DistServe: Disaggregating prefill and decoding for goodput-optimised LLM serving,” *OSDI*, 2024. <https://arxiv.org/abs/2401.09670>
- [8] MLCommons, “MLPerf Inference v5.1 small-LLM track,” <https://mlcommons.org/2025/09/small-llm-inference-5-1/>, 2025.
- [9] NVIDIA Developer Blog, “LLM benchmarking fundamental concepts,” <https://developer.nvidia.com/blog/llm-benchmarking-fundamental-concepts/>.
- [10] Anyscale, “LLM serving benchmarking metrics,” <https://docs.anyscale.com/llm/serving/benchmarking/metrics>.
- [11] Z. Li *et al.*, “prima.cpp: piped-ring parallel inference for 70B-class LLMs on home clusters,” *arXiv:2504.08791*, 2025. <https://arxiv.org/abs/2504.08791>
- [12] A. Q. Jiang *et al.*, “Mixtral of Experts,” *arXiv:2401.04088*, 2024.
- [13] Qwen Team, “Qwen3-30B-A3B-Instruct model card,” Hugging Face, 2025. <https://huggingface.co/Qwen/Qwen3-30B-A3B-Instruct>
- [14] exo-explore, “EXO — run frontier AI locally,” <https://github.com/exo-explore/exo>.
- [15] P. Zheng, W. Xu, H. Wang, J. Chen, X. Shen, “HALO: Semantic-aware distributed LLM inference in lossy edge network,” *arXiv:2601.11676*, January 2026. <https://arxiv.org/abs/2601.11676>
- [16] Groq, “Introducing Llama-3-Groq-Tool-Use Models,” 2024. <https://groq.com/blog/introducing-llama-3-groq-tool-use-models>
- [17] N. Muennighoff *et al.*, “OLMoE: Open Mixture-of-Experts Language Models,” *arXiv:2409.02060*, 2024. <https://huggingface.co/allenai/OLMoE-1B-7B-0924>
- [18] PyTorch / Meta, “Introducing Async Tensor Parallelism in PyTorch,” *PyTorch Discuss forums (TorchTitan project)*, September 2024. <https://discuss.pytorch.org/t/distributed-w-torchtitan-introducing-async-tensor-parallelism-in-pytorch/209487>
- [19] L. Gao *et al.*, “A framework for few-shot language model evaluation (`lm-evaluation-harness`),” EleutherAI, v0.4.x, 2024–2026. <https://github.com/EleutherAI/lm-evaluation-harness>
- [20] ByteShape, “A 30B Qwen model walks into a Raspberry Pi... and runs in real time,” technical blog, 2026. <https://byteshape.com/blogs/Qwen3-30B-A3B-Instruct-2507/>

- [21] J. Geerling, “Qwen3-30B-A3B on Raspberry Pi 5 16GB,” `geerlingguy/ai-benchmarks` issue #47, 2026.
<https://github.com/geerlingguy/ai-benchmarks/issues/47>

*Document generated 23 May 2026 from the operational cluster. HELLOMATIK.COM · Raspberry Pi Cluster Lab.
Final result 15.143 tok/s decode, +16.1% above the public ceiling (14.449 tok/s end-to-end sustained).*